

mszcuxc3a

March 20, 2025

```
[15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from sklearn.metrics import mean_absolute_error, r2_score
```

```
[33]: # Importing DataSet and take a look at Data
data = pd.read_csv('housingdata.csv')
data
```

```
[33]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	NaN	36.2
..
501	21.0	391.99	NaN	22.4
502	21.0	396.90	9.08	20.6

```
503      21.0  396.90   5.64  23.9
504      21.0  393.45   6.48  22.0
505      21.0  396.90   7.88  11.9
```

```
[506 rows x 14 columns]
```

```
[34]: data.isnull().sum()
```

```
[34]: CRIM      20
      ZN       20
      INDUS   20
      CHAS    20
      NOX      0
      RM       0
      AGE     20
      DIS      0
      RAD      0
      TAX      0
      PTRATIO  0
      B        0
      LSTAT   20
      MEDV     0
      dtype: int64
```

```
[35]: # Handle null values by filling them with the mean of the respective columns
      data.fillna(data.mean(), inplace=True)
```

```
[36]: data.isnull().sum()
```

```
[36]: CRIM      0
      ZN       0
      INDUS   0
      CHAS    0
      NOX      0
      RM       0
      AGE     0
      DIS      0
      RAD      0
      TAX      0
      PTRATIO  0
      B        0
      LSTAT   0
      MEDV     0
      dtype: int64
```

```
[37]: data.describe()
```

```
[37]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.611874	11.211934	11.083992	0.069959	0.554695	6.284634
std	8.545770	22.921051	6.699165	0.250233	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.083235	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.290250	0.000000	9.900000	0.000000	0.538000	6.208500
75%	3.611874	11.211934	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.518519	3.795043	9.549407	408.237154	18.455534	356.674032
std	27.439466	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.925000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	74.450000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	93.575000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.715432	22.532806
std	7.012739	9.197104
min	1.730000	5.000000
25%	7.230000	17.025000
50%	11.995000	21.200000
75%	16.570000	25.000000
max	37.970000	50.000000


```
[38]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    int64
9   TAX         506 non-null    int64
```

```
10 PTRATIO  506 non-null    float64
11 B        506 non-null    float64
12 LSTAT    506 non-null    float64
13 MEDV     506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```


```
[39]: import seaborn as sns
      sns.distplot(data.MEDV)
```

```
[39]: <Axes: xlabel='MEDV', ylabel='Density'>
```



```
[40]: sns.boxplot(data.MEDV)
```

```
[40]: <Axes: ylabel='MEDV'>
```



```
[41]: correlation = data.corr()
      correlation.loc['MEDV']
```

```
[41]: CRIM      -0.379695
      ZN        0.365943
      INDUS    -0.478657
      CHAS      0.179882
      NOX      -0.427321
      RM        0.695360
      AGE      -0.380223
      DIS       0.249929
      RAD      -0.381626
      TAX      -0.468536
      PTRATIO  -0.507787
      B         0.333461
      LSTAT    -0.721975
      MEDV      1.000000
      Name: MEDV, dtype: float64
```


```
[42]: # plotting the heatmap
      import matplotlib.pyplot as plt
      fig, axes = plt.subplots(figsize=(15,12))
      sns.heatmap(correlation, square = True, annot = True)
```

[42]: <Axes: >



[43]: # Checking the scatter plot with the most correlated features

```
plt.figure(figsize = (20,5))
features = ['LSTAT','RM','PTRATIO']
for i, col in enumerate(features):
    plt.subplot(1, len(features) , i+1)
    x = data[col]
    y = data.MEDV
    plt.scatter(x, y, marker='o')
    plt.title("Variation in House prices")
    plt.xlabel(col)
    plt.ylabel('"House prices in $1000"')
```



```
[44]: # Splitting the dependent feature and independent feature
#X = data[['LSTAT', 'RM', 'PTRATIO']]
X = data.iloc[:, :-1]
y= data.MEDV
```

```
[45]: import numpy as np
from sklearn.model_selection import train_test_split

# Assuming you have data stored in some variables X and y
# Splitting data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Now you can proceed with the code you provided
# Importing necessary libraries
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler

# Scaling the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

mean = X_train.mean(axis=0)
std = X_train.std(axis=0)
X_train = (X_train - mean) / std
X_test = (X_test - mean) / std
#Linear Regression

from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
#Fitting the model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
regressor.fit(X_train,y_train)
```

```
# Model Evaluation
```

```
[45]: LinearRegression()
```

```
[46]: #Prediction on the test dataset  
y_pred = regressor.predict(X_test)  
# Predicting RMSE the Test set results  
from sklearn.metrics import mean_squared_error  
rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))  
print(rmse)
```

```
5.001766890194174
```

```
[47]: from sklearn.metrics import r2_score  
r2 = r2_score(y_test, y_pred)  
print(r2)
```

```
0.658852019550814
```

```
[48]: import keras  
from keras.layers import Dense  
from keras.models import Sequential  
from sklearn.preprocessing import StandardScaler  
import matplotlib.pyplot as plt  
  
# Assuming X_train and X_test are defined and initialized previously  
# Assuming y_train is also defined and initialized  
  
# Scaling the dataset  
sc = StandardScaler()  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)  
  
# Creating the neural network model  
model = Sequential()  
model.add(Dense(128, activation='relu', input_dim=13))  
model.add(Dense(64, activation='relu'))  
model.add(Dense(32, activation='relu'))  
model.add(Dense(16, activation='relu'))  
model.add(Dense(1))  
  
# Compiling the model  
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])  
  
# Visualizing the model architecture  
keras.utils.plot_model(model, to_file='model.png', show_shapes=True,  
    ↪ show_layer_names=True)
```



```

# Assuming you have defined your training data X_train and y_train
history = model.fit(X_train, y_train, epochs=100, validation_split=0.05)

# Plotting the training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

You must install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model` to work.

```

Epoch 1/100
12/12          6s 63ms/step -
loss: 637.5823 - mae: 23.3392 - val_loss: 501.9428 - val_mae: 20.9218
Epoch 2/100
12/12          0s 16ms/step -
loss: 576.0250 - mae: 21.9062 - val_loss: 426.7097 - val_mae: 19.0900
Epoch 3/100
12/12          0s 16ms/step -
loss: 472.9474 - mae: 19.5632 - val_loss: 260.3259 - val_mae: 14.2853
Epoch 4/100
12/12          0s 22ms/step -
loss: 226.6125 - mae: 12.8536 - val_loss: 70.4344 - val_mae: 5.7795
Epoch 5/100
12/12          1s 18ms/step -
loss: 101.2044 - mae: 7.9105 - val_loss: 63.4414 - val_mae: 5.2930
Epoch 6/100
12/12          0s 18ms/step -
loss: 66.6758 - mae: 6.2065 - val_loss: 67.2448 - val_mae: 5.1342
Epoch 7/100
12/12          0s 19ms/step -
loss: 39.7218 - mae: 4.6565 - val_loss: 56.6962 - val_mae: 4.7042
Epoch 8/100
12/12          0s 16ms/step -
loss: 22.6439 - mae: 3.5725 - val_loss: 53.8994 - val_mae: 4.7284
Epoch 9/100
12/12          0s 17ms/step -
loss: 22.5735 - mae: 3.4815 - val_loss: 55.6253 - val_mae: 4.6438
Epoch 10/100
12/12          0s 19ms/step -
loss: 20.3210 - mae: 3.1774 - val_loss: 55.0472 - val_mae: 4.5780
Epoch 11/100
12/12          0s 22ms/step -
loss: 19.6458 - mae: 3.1178 - val_loss: 53.4736 - val_mae: 4.5105

```

Epoch 12/100
12/12 0s 16ms/step -
loss: 15.4462 - mae: 2.8748 - val_loss: 50.7672 - val_mae: 4.3458
Epoch 13/100
12/12 0s 15ms/step -
loss: 19.8370 - mae: 3.2098 - val_loss: 49.7969 - val_mae: 4.3428
Epoch 14/100
12/12 0s 18ms/step -
loss: 14.6462 - mae: 2.7288 - val_loss: 49.2393 - val_mae: 4.3534
Epoch 15/100
12/12 0s 13ms/step -
loss: 15.0232 - mae: 2.9019 - val_loss: 45.9021 - val_mae: 4.2246
Epoch 16/100
12/12 0s 11ms/step -
loss: 15.3578 - mae: 2.8083 - val_loss: 46.2940 - val_mae: 4.2638
Epoch 17/100
12/12 0s 18ms/step -
loss: 13.5767 - mae: 2.6155 - val_loss: 44.7240 - val_mae: 4.1594
Epoch 18/100
12/12 0s 17ms/step -
loss: 12.9807 - mae: 2.6075 - val_loss: 44.4587 - val_mae: 4.1081
Epoch 19/100
12/12 0s 19ms/step -
loss: 12.7454 - mae: 2.6092 - val_loss: 44.1038 - val_mae: 4.1315
Epoch 20/100
12/12 0s 18ms/step -
loss: 10.9545 - mae: 2.4212 - val_loss: 42.1538 - val_mae: 4.0154
Epoch 21/100
12/12 0s 18ms/step -
loss: 13.6765 - mae: 2.6469 - val_loss: 42.3075 - val_mae: 4.0334
Epoch 22/100
12/12 0s 14ms/step -
loss: 11.0072 - mae: 2.4324 - val_loss: 42.4709 - val_mae: 4.0275
Epoch 23/100
12/12 0s 15ms/step -
loss: 11.9222 - mae: 2.4238 - val_loss: 40.3017 - val_mae: 3.9351
Epoch 24/100
12/12 0s 13ms/step -
loss: 11.6086 - mae: 2.5241 - val_loss: 41.2279 - val_mae: 4.0053
Epoch 25/100
12/12 0s 28ms/step -
loss: 10.6481 - mae: 2.3478 - val_loss: 40.2435 - val_mae: 3.9149
Epoch 26/100
12/12 0s 18ms/step -
loss: 13.4676 - mae: 2.6480 - val_loss: 39.4666 - val_mae: 3.9127
Epoch 27/100
12/12 0s 15ms/step -
loss: 11.2125 - mae: 2.4112 - val_loss: 39.4673 - val_mae: 3.8951


Epoch 28/100
12/12 0s 16ms/step -
loss: 11.7875 - mae: 2.5431 - val_loss: 36.6572 - val_mae: 3.7522
Epoch 29/100
12/12 0s 16ms/step -
loss: 11.2983 - mae: 2.4388 - val_loss: 38.5011 - val_mae: 3.9132
Epoch 30/100
12/12 0s 17ms/step -
loss: 11.5294 - mae: 2.3625 - val_loss: 38.0044 - val_mae: 3.8109
Epoch 31/100
12/12 0s 19ms/step -
loss: 10.8664 - mae: 2.4427 - val_loss: 39.1625 - val_mae: 3.9447
Epoch 32/100
12/12 0s 11ms/step -
loss: 11.4726 - mae: 2.4421 - val_loss: 37.6246 - val_mae: 3.8349
Epoch 33/100
12/12 0s 15ms/step -
loss: 10.3745 - mae: 2.3295 - val_loss: 36.4306 - val_mae: 3.7920
Epoch 34/100
12/12 0s 15ms/step -
loss: 8.2271 - mae: 2.0752 - val_loss: 37.1040 - val_mae: 3.8626
Epoch 35/100
12/12 0s 16ms/step -
loss: 7.5315 - mae: 1.9786 - val_loss: 36.1065 - val_mae: 3.7770
Epoch 36/100
12/12 0s 16ms/step -
loss: 9.0459 - mae: 2.1542 - val_loss: 37.0387 - val_mae: 3.8579
Epoch 37/100
12/12 0s 15ms/step -
loss: 8.7370 - mae: 2.1251 - val_loss: 35.5475 - val_mae: 3.7793
Epoch 38/100
12/12 0s 18ms/step -
loss: 9.1603 - mae: 2.2012 - val_loss: 34.4235 - val_mae: 3.7528
Epoch 39/100
12/12 0s 15ms/step -
loss: 9.5289 - mae: 2.2026 - val_loss: 34.4647 - val_mae: 3.7543
Epoch 40/100
12/12 0s 14ms/step -
loss: 8.6058 - mae: 2.1994 - val_loss: 34.3443 - val_mae: 3.7876
Epoch 41/100
12/12 0s 16ms/step -
loss: 8.6005 - mae: 2.1878 - val_loss: 34.6477 - val_mae: 3.7994
Epoch 42/100
12/12 0s 15ms/step -
loss: 8.1823 - mae: 2.1236 - val_loss: 35.1218 - val_mae: 3.7617
Epoch 43/100
12/12 0s 17ms/step -
loss: 8.4245 - mae: 2.1372 - val_loss: 34.2113 - val_mae: 3.7304

Epoch 44/100
12/12 0s 16ms/step -
loss: 8.1443 - mae: 2.1329 - val_loss: 34.0813 - val_mae: 3.6969
Epoch 45/100
12/12 0s 21ms/step -
loss: 7.4777 - mae: 2.0959 - val_loss: 31.5419 - val_mae: 3.6159
Epoch 46/100
12/12 0s 27ms/step -
loss: 8.7502 - mae: 2.1694 - val_loss: 32.1069 - val_mae: 3.6619
Epoch 47/100
12/12 0s 16ms/step -
loss: 7.2737 - mae: 2.0741 - val_loss: 31.9654 - val_mae: 3.6828
Epoch 48/100
12/12 0s 34ms/step -
loss: 7.6746 - mae: 2.1023 - val_loss: 31.7952 - val_mae: 3.6014
Epoch 49/100
12/12 1s 22ms/step -
loss: 7.0337 - mae: 1.9327 - val_loss: 31.1338 - val_mae: 3.5977
Epoch 50/100
12/12 0s 24ms/step -
loss: 7.3431 - mae: 2.0354 - val_loss: 31.8056 - val_mae: 3.6743
Epoch 51/100
12/12 1s 76ms/step -
loss: 7.8528 - mae: 2.0837 - val_loss: 33.4314 - val_mae: 3.7830
Epoch 52/100
12/12 1s 27ms/step -
loss: 7.5038 - mae: 1.9968 - val_loss: 29.3334 - val_mae: 3.4450
Epoch 53/100
12/12 1s 48ms/step -
loss: 8.9781 - mae: 2.2202 - val_loss: 32.2382 - val_mae: 3.6970
Epoch 54/100
12/12 0s 21ms/step -
loss: 7.4570 - mae: 2.0201 - val_loss: 28.3175 - val_mae: 3.4512
Epoch 55/100
12/12 0s 20ms/step -
loss: 7.5158 - mae: 2.0010 - val_loss: 30.5191 - val_mae: 3.6869
Epoch 56/100
12/12 0s 20ms/step -
loss: 7.1159 - mae: 2.0444 - val_loss: 27.9140 - val_mae: 3.4476
Epoch 57/100
12/12 0s 21ms/step -
loss: 6.5534 - mae: 1.9562 - val_loss: 29.9596 - val_mae: 3.5577
Epoch 58/100
12/12 0s 20ms/step -
loss: 7.2865 - mae: 2.0860 - val_loss: 31.7106 - val_mae: 3.7737
Epoch 59/100
12/12 0s 23ms/step -
loss: 6.8538 - mae: 1.9976 - val_loss: 28.0757 - val_mae: 3.4892

Epoch 60/100
12/12 0s 20ms/step -
loss: 7.3235 - mae: 1.9828 - val_loss: 26.4554 - val_mae: 3.4103
Epoch 61/100
12/12 0s 19ms/step -
loss: 5.7066 - mae: 1.8667 - val_loss: 27.6188 - val_mae: 3.4808
Epoch 62/100
12/12 0s 18ms/step -
loss: 6.8279 - mae: 1.9656 - val_loss: 29.8978 - val_mae: 3.6490
Epoch 63/100
12/12 0s 19ms/step -
loss: 6.9390 - mae: 1.9441 - val_loss: 26.4475 - val_mae: 3.3885
Epoch 64/100
12/12 0s 23ms/step -
loss: 6.1152 - mae: 1.9217 - val_loss: 26.6189 - val_mae: 3.4163
Epoch 65/100
12/12 0s 21ms/step -
loss: 5.8128 - mae: 1.8767 - val_loss: 28.0359 - val_mae: 3.5491
Epoch 66/100
12/12 0s 23ms/step -
loss: 6.2168 - mae: 1.9205 - val_loss: 26.7844 - val_mae: 3.4631
Epoch 67/100
12/12 0s 16ms/step -
loss: 6.2298 - mae: 1.8120 - val_loss: 27.2993 - val_mae: 3.4859
Epoch 68/100
12/12 1s 43ms/step -
loss: 5.8782 - mae: 1.8068 - val_loss: 26.5276 - val_mae: 3.4292
Epoch 69/100
12/12 0s 20ms/step -
loss: 6.1733 - mae: 1.8634 - val_loss: 28.6765 - val_mae: 3.6696
Epoch 70/100
12/12 0s 18ms/step -
loss: 4.8117 - mae: 1.6656 - val_loss: 23.6609 - val_mae: 3.2500
Epoch 71/100
12/12 0s 17ms/step -
loss: 5.3697 - mae: 1.7578 - val_loss: 26.0816 - val_mae: 3.4494
Epoch 72/100
12/12 0s 16ms/step -
loss: 4.8997 - mae: 1.6800 - val_loss: 24.6617 - val_mae: 3.3255
Epoch 73/100
12/12 0s 16ms/step -
loss: 5.0019 - mae: 1.7287 - val_loss: 24.3470 - val_mae: 3.3590
Epoch 74/100
12/12 0s 22ms/step -
loss: 4.9347 - mae: 1.6787 - val_loss: 23.5476 - val_mae: 3.3513
Epoch 75/100
12/12 0s 19ms/step -
loss: 5.2418 - mae: 1.7595 - val_loss: 24.8325 - val_mae: 3.3322

Epoch 76/100
12/12 0s 25ms/step -
loss: 5.4352 - mae: 1.7829 - val_loss: 24.2416 - val_mae: 3.3624
Epoch 77/100
12/12 0s 21ms/step -
loss: 5.2187 - mae: 1.7652 - val_loss: 25.5319 - val_mae: 3.4924
Epoch 78/100
12/12 0s 21ms/step -
loss: 5.3340 - mae: 1.7537 - val_loss: 25.4600 - val_mae: 3.4550
Epoch 79/100
12/12 0s 19ms/step -
loss: 5.2333 - mae: 1.7186 - val_loss: 24.4441 - val_mae: 3.3485
Epoch 80/100
12/12 0s 23ms/step -
loss: 4.5731 - mae: 1.5979 - val_loss: 23.2122 - val_mae: 3.3367
Epoch 81/100
12/12 1s 48ms/step -
loss: 5.5938 - mae: 1.8376 - val_loss: 25.4527 - val_mae: 3.4370
Epoch 82/100
12/12 1s 22ms/step -
loss: 4.3674 - mae: 1.6225 - val_loss: 24.8968 - val_mae: 3.4075
Epoch 83/100
12/12 0s 15ms/step -
loss: 5.1668 - mae: 1.7162 - val_loss: 24.7691 - val_mae: 3.3807
Epoch 84/100
12/12 0s 14ms/step -
loss: 4.8396 - mae: 1.6666 - val_loss: 23.3055 - val_mae: 3.3047
Epoch 85/100
12/12 0s 15ms/step -
loss: 4.9508 - mae: 1.7420 - val_loss: 22.3698 - val_mae: 3.1932
Epoch 86/100
12/12 0s 21ms/step -
loss: 5.2578 - mae: 1.7667 - val_loss: 23.5968 - val_mae: 3.3787
Epoch 87/100
12/12 0s 16ms/step -
loss: 5.3156 - mae: 1.6898 - val_loss: 23.8333 - val_mae: 3.3489
Epoch 88/100
12/12 0s 23ms/step -
loss: 4.4919 - mae: 1.6335 - val_loss: 21.3748 - val_mae: 3.1345
Epoch 89/100
12/12 0s 21ms/step -
loss: 5.1205 - mae: 1.7094 - val_loss: 23.9714 - val_mae: 3.3762
Epoch 90/100
12/12 0s 21ms/step -
loss: 4.6511 - mae: 1.6098 - val_loss: 20.4697 - val_mae: 3.1307
Epoch 91/100
12/12 1s 40ms/step -
loss: 3.7245 - mae: 1.4520 - val_loss: 21.4718 - val_mae: 3.1953

Epoch 92/100
12/12 0s 21ms/step -
loss: 4.3476 - mae: 1.5853 - val_loss: 22.9994 - val_mae: 3.2758
Epoch 93/100
12/12 0s 23ms/step -
loss: 4.6215 - mae: 1.5950 - val_loss: 21.7747 - val_mae: 3.1819
Epoch 94/100
12/12 0s 25ms/step -
loss: 4.1278 - mae: 1.5520 - val_loss: 22.8728 - val_mae: 3.3002
Epoch 95/100
12/12 1s 37ms/step -
loss: 4.2500 - mae: 1.5774 - val_loss: 20.0942 - val_mae: 3.0648
Epoch 96/100
12/12 0s 22ms/step -
loss: 4.2548 - mae: 1.5868 - val_loss: 21.3404 - val_mae: 3.1491
Epoch 97/100
12/12 0s 27ms/step -
loss: 4.4553 - mae: 1.6012 - val_loss: 23.8823 - val_mae: 3.3865
Epoch 98/100
12/12 0s 26ms/step -
loss: 4.1180 - mae: 1.4955 - val_loss: 20.3904 - val_mae: 3.0532
Epoch 99/100
12/12 1s 16ms/step -
loss: 4.2431 - mae: 1.5470 - val_loss: 20.9968 - val_mae: 3.1629
Epoch 100/100
12/12 0s 27ms/step -
loss: 4.5520 - mae: 1.6360 - val_loss: 21.0372 - val_mae: 3.1553



```
[49]: # Plotting the Mean Absolute Error using Plotly
import plotly.graph_objects as go
fig = go.Figure()
fig.add_trace(go.Scattergl(y=history.history['mae'], name='Train'))
fig.add_trace(go.Scattergl(y=history.history['val_mae'], name='Valid'))
fig.update_layout(height=500, width=700, xaxis_title='Epoch', yaxis_title='Mean_
↪Absolute Error')
fig.show()
```

```
[50]: #Evaluation of the model
y_pred = model.predict(X_test)
mse_nn, mae_nn = model.evaluate(X_test, y_test)
print('Mean squared error on test data: ', mse_nn)
print('Mean absolute error on test data: ', mae_nn)
```

```
4/4          2s 241ms/step
4/4          0s 25ms/step - loss:
9.5879 - mae: 2.1784
Mean squared error on test data: 12.936566352844238
Mean absolute error on test data: 2.3468189239501953
```



```
[51]: #Comparison with traditional approaches
      #First let's try with a simple algorithm, the Linear Regression:
      from sklearn.metrics import mean_absolute_error
      lr_model = LinearRegression()
      lr_model.fit(X_train, y_train)
      y_pred_lr = lr_model.predict(X_test)
      mse_lr = mean_squared_error(y_test, y_pred_lr)
      mae_lr = mean_absolute_error(y_test, y_pred_lr)
      print('Mean squared error on test data: ', mse_lr)
      print('Mean absolute error on test data: ', mae_lr)
      from sklearn.metrics import r2_score
      r2 = r2_score(y_test, y_pred)
      print(r2)
```

Mean squared error on test data: 25.017672023842856
Mean absolute error on test data: 3.1499233573458025
0.8235933681883683

```
[52]: # Predicting RMSE the Test set results
      from sklearn.metrics import mean_squared_error
      rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
      print(rmse)
```

3.5967437614204782

```
[53]: # Make predictions on new data
      import sklearn
      new_data = sklearn.preprocessing.StandardScaler().fit_transform([[0.1, 10.0,
      5.0, 0, 0.4, 6.0, 50, 6.0, 1, 400, 20, 300, 10]])
      prediction = model.predict(new_data)
      print("Predicted house price:", prediction)
```

1/1 0s 484ms/step
Predicted house price: [[11.963507]]